

Ash FS: A Flash Drive File System

Daniel Băluță

Faculty of Computer Science and Engineering
Politehnica University of Bucharest
Romania, Bucharest
Email: daniel.baluta@gmail.com

Gabriel Sandu

Faculty of Computer Science and Engineering
Politehnica University of Bucharest
Romania, Bucharest
Email: gabrim.san@gmail.com

Abstract—The recent increase in USB Flash Drive capacities has triggered a need for a special file system that can handle large files such as the ones encountered in a NTFS environment, while keeping a good error correction and access speed.

In Ash FS, we are trying to reduce the wear levelling of the flash drive by storing such large quantity of data in an uniform way and provide the possibility of using compression, crypting and error correction techniques.

I. INTRODUCTION

A. Flash

Flash memory is non-volatile computer memory that can be electrically erased and reprogramed. During last years it become a common storage medium in embedded devices, because it provides solid state storage with high reliability, at a relatively low cost.

Flash is a specific type of EEPROM (Electrically Erasable Programmable Read-Only Memory) that is erased and programmed in large blocks; in early flash the entire chip had to be erased at once. It is available in two major types - the traditional NOR flash which is directly accessible, and the newer, cheaper NAND flash which is addressable only through a single 8-bit bus used for both data and addresses, with separate control lines. This types of flash share their most important characteristics - each bit in a clean flash chip will be set to a logical one, and can be set to zero by a write operation.

Flash chips are arranged into blocks which are typically 128KB on NOR flash and 8KB on NAND flash. Resetting bits from zero to one cannot be done individually, but only by resetting or erasing a complete block. The lifetime of a flash chip is measured in such erase cycles, with the typical lifetime being 100,000 erases per block. To ensure that no erase block reaches this limit before the rest of the chip, most users of flash chips attempt to ensure that erase cycles are evenly distributed around the flash; process known as "wear levelling".

Aside from the difference in erase block sizes, NAND flash chips also have other differences from NOR chips. They are further divided into "pages" which are typically 512 bytes

in size, each of which has an extra 16 bytes of "out of band" storage space, intended to be used for metadata or error correction. NAND flash is written by loading the required data into an internal buffer one byte at a time, then issuing a write command. While NOR flash allows bits to be cleared individually until there are none left to be cleared, NAND flash allows only ten such write cycles to each page before leakage causes contents to become undefined until the next erase of the block in which the page resides.

B. Flash Translation Layers

The majority of applications of flash for file storage have involved using the flash to emulate a block device with standard 512-byte sectors, and then using standard file systems on that emulated device.

The simplest method of achieving this is to use a simple 1:1 mapping from the emulated block device to the flash chip, and to simulate the smaller sector size for write requests by reading the whole erase block, modifying the appropriate part of the buffer, erasing and rewriting the entire block. This approach provides no wear levelling, and is extremely unsafe because of the potential power loss between the erase subsequent rewrite of the data. However, it is acceptable for use during development of a file system which is intended for read-only operation in production modes. The mtdblock Linux driver provides this functionality, slightly optimised to prevent excessive erase cycles by gathering writes to a single erase block and only performing the erase/modify/writeback procedure when a write to a different erase block is requested.

To emulate a block device in a fashion suitable for use with a writable file system, a more sophisticated approach is required.

To provide wear levelling and reliable operation, sectors of the emulated block device, are stored in varying locations on the physical medium, and a "Translation Layer" is used to keep track of current location of each sector in the emulated block device. This translation layers is effectively a form of journalling file system.

The most common such translation layer is a component of PCMCIA specification, the "Flash Translation Layer"

(FTL). More recently, a variant designed for use with NAND flash chips has been in widespread use in the popular DiskOnChip devices produced by M-Systems.

Unfortunately, both FTL and the newer NFTL are encumbered by patents. M-Systems have granted a license for FTL to be used on all PCMCIA devices, and allow NFTL to be used only on DiskOnChip devices.

Linux supports both of these translation layers, but their use is deprecated and intended for backwards compatibility only. Not only are there patent issues, but the practice of using a form of journalling file system to emulate a block device, on which a "standard" journalling file system is then used, is unnecessarily inefficient.

A far more efficient use of flash technology would be permitted by the use of a file system designed specifically for use on such devices, with no extra layers of translation in between.

Ash File System is a block device based file system and tries to take in consideration the particular characteristics of flash memory.

II. VIRTUAL FILE SYSTEM

Before diving into the Ash File System we must have a short look at the Virtual File System (VFS). A virtual file system is an abstraction layer on the top of a more concrete file system and supports a uniform view of the objects or files in the filesystem. Even though the meaning of the term file may appear to be clear, there are many small, often subtle differences in detail owing to the underlying implementations of the individual filesystems. Not all support the same functions, and some operations make no sense when applied to certain objects (e.g. named pipes).

When working with files, the central objects differ in kernel space and userspace. For user programs, a file is identified by a file descriptor. This is an integer number used as a parameter to identify the file in all file-related operations. The file descriptor is assigned by the kernel when a file is opened and is valid only within a process. Two different processes may therefore use the same file descriptor, but it does not point to the same file in both cases. Shared use of files on the basis of the same descriptor number is not possible.

The inode is key to the kernel's work with files. Each file (and each directory) has just one inode, which contains metadata such as access rights, date of last change, and so on, and also pointers to the file data. However, the inode does not contain one important item of information - the filename. Usually, it is assumed that the name of the file is one of its major characteristics and should therefore be included in the object (inode) used to manage it.

A. Structure of the VFS

The key idea behind the VFS consists of introducing a common file model capable of representing all supported filesystems. This model strictly follows the file model provided by the traditional Unix filesystem. However, each specific filesystem implementation must translate its physical organization into the VFS's common file model. For instance, in the common file model, each directory is regarded as a file, which contains a list of files and other directories. Several non-Unix disk-based filesystems use a File Allocation Table (FAT), which stores the position of each file in the directory tree. In these filesystems, directories are not files. To stick to the VFS's common file model, the Linux implementation of such FAT-based filesystems must be able to construct on the fly, when needed, the files corresponding to the directories. Such files only exist only as objects in kernel memory.

Let's see an example that illustrates this concept by showing how the `read()` operation would be translated by the kernel into a call specific to the MS-DOS filesystem. The userspace application's call to `read()` makes the kernel invoke the corresponding `sys_read` function, like every other system call. The file is represented, in kernel, by a file data structure which contains a field called `f_op` that has pointers to functions specific to MS-DOS files, including a function that reads a file. `sys_read()` finds the pointer to this function and invokes it. Thus, the application's `read()` is turned into the rather indirect call: `file.f_op.read(..)`.

The common filesystem model consists of the following object types:

- **superblock object**, stores information concerning a mounted filesystem. For disk-based filesystems, this object usually corresponds to a filesystem control block stored on disk.
- **inode object**, stores general information about a specific file. For disk-based filesystems, this object usually corresponds to a file control block stored on disk. Each inode object is associated with an inode number, which uniquely identifies the file within the filesystem.
- **file object**, stores information about the interaction between an open file and a process. This information exists only in kernel memory during the period when a process has the file open.
- **dentry object**, stores information about the linking of a directory entry (that is, a particular name of the file) with the corresponding file. Each disk-based filesystem stores this information in its own particular way on disk.

Besides providing a common interface to all filesystem implementations, the VFS has another important role related to system performance. The most recently used dentry objects are contained in a disk cache named the dentry cache, which speeds up the translation from a file pathname to the inode of the last pathname component.

B. VFS Data Structures

Now let's take a look at a detailed description of filesystem's objects:

1) *The superblock*: The superblock object is implemented by each filesystem, and is used to store information describing the specific filesystem. This object usually corresponds to the filesystem superblock or the filesystem control block, which is stored in a special sector on disk. The superblock is read in memory at filesystem mount time.

The superblock object is represented by *struct super_block* and defined in *linux/fs.h*. The code for creating, managing and destroying superblock objects can be found in *fs/super.c*. A superblock object is created with the *alloc_super()* function. When mounted, a filesystem invokes this function, reads its superblock off the disk, and fills in its superblock object.

The most important data member in the superblock object is *s_op*, which is the superblock operations table. The superblock operations table is represented by *struct super_operations* and is defined in *fs.h* and it contains pointers to functions operating on a superblock object. When a filesystem needs to perform an operation on its superblock, it follows the pointers from its superblock object to the desired method. For example if a filesystem wanted to write to its superblock, it would invoke

sb → *s_op* → *write_super(sb)*

where *sb* is a pointer to the filesystem's superblock. Following that pointer into *s_op* yields the superblock operations table and ultimately the desired *write_super*. It is beyond the scope of this paper to discuss all superblock operations but its worth pointing that a specific filesystem can set one or more superblock operations to *NULL* in which case the VFS either calls a generic function or does nothing depending on the operation.

2) *The inode*: The inode object represents all the information needed by the kernel to manipulate a file or a directory. If a filesystem does not have inodes, the filesystem must obtain the information from wherever it is stored (e.g. inode is embedded into the files). The inode object is represented by *struct inode* and defined in *linux/fs.h* and it is constructed in memory only when the files are accessed. Some of the entries in *struct inode* are related to these special files. For example, *i_pipe* field points to a named pipe data structure. If the inode does not refer to a named pipe, this field is set to *NULL*.

As with the superblock operations, the *inode_operations* member is very important. It describes the filesystem's implemented functions that the VFS invoke on an inode and their definitions can be found in *linux/fs.h* file.

3) *The file*: The file object is used to represent a file opened by a process. When we think of the VFS from the perspective of user-space, the file object is what readily comes to mind. Processes deal directly with files, not superblocks,

inodes, or dentries. It is not surprising that the information in the file object is the most familiar (data such as access mode and current offset) or that the file operations are familiar system calls such as *read()* and *write()*.

The file object is the in-memory representation of an open file. The object (but not the physical file) is created in response to the *open()* system call and destroyed in response to the *close()* system call. All these file-related calls are actually methods defined in the file operations table. Because multiple processes can open and manipulate a file at the same time, there can be multiple file objects in existence for the same file. This is why when managing file pointer every process has its own view on the read/write position. The file object merely represents a process's view of an open file. The object points back to the dentry (which in turn points back to the inode) that actually represents the open file. The inode and dentry objects, of course, are unique.

The file object is represented by *struct file* and it is defined in *linux/fs.h*. The file object does not actually correspond to any on-disk data. Therefore, there is no flag in the object to represent whether the object is dirty and needs to be written back to disk. The file object points to its associated dentry object using *f_dentry* pointer. The dentry in turn points to the associated inode, which reflects whether the file is dirty.

As with all the other VFS objects, the file operations table is very important. The operations associated with *struct file* are the familiar system calls that form the basis of the Unix system calls. The file object methods are specified in *file_operations* and defined in *linux/fs.h*.

4) *The dentry*: The Virtual File System treats directories as files. An inode object represents both these components. Despite this useful unification, the VFS often needs to perform directory-specific operations, such as path name lookup. Path name lookup involves translating each component of a path, ensuring it is valid, and following it to the next component. To facilitate this, the VFS employs the concept of a directory entry (dentry). A dentry is a specific component in a path. Dentry objects are all components in a path, including files. Resolving a path and walking its components is a nontrivial exercise, time-consuming and rife with string comparisons. The dentry object makes the whole process easier. Dentries might also include mount points. The VFS constructs dentry objects on the fly, as needed, when performing directory operations.

Dentry objects are represented by *struct dentry* and defined in *linux/dcache.h* and it does not correspond to any sort of on-disk data structure. The VFS creates the dentry from the string representation of a path name. Because the dentry object is not physically stored on the disk, no flag in *struct dentry* specifies whether the object is dirty or not.

The *dentry_operations* structure specifies the methods that



Fig. 1. Write Speed Test




Fig. 2. Read Speed Test




Fig. 3. Crypting and Compression Speed Test

virtual file system invokes on directory entries on a given filesystem. The *dentry_operations* structure is defined in *linux/dcache.h*

III. ASH FILE SYSTEM

A. Storage Format

Add text for Storage Format

B. Mounting

Add text for Mounting

C. Operations

Add text for operations

D. Cryptography

Add text for cryptography

E. Compression

Add text for compression

IV. FUTURE WORK

A. Combine compression with encryption

In the current version both compression and encryption are supported but separate. At any time you can have enabled either compression or encryption but not both in the same time. We must analyse the possibility of integrating compression with encryption and see the possible benefits. At a first glance there is one question arising - which of them should be done first.

B. eXecute In Place

One very nice and useful feature would be *eXecute In Place* (XIP) which assumes that code is directly executed from the flash drive. When a program residing on *AshFS* is run , the executable code is copied from flash into RAM before the CPU can execute it. Even , when the *mmap()* system call is implied data are not accessed directly from the flash but are copied into RAM.

It is clear that, XIP and compression are mutually exclusive, if data are compressed they cannot be used directly in place. It is interesting to see then, what is the optimal solution: save amount of RAM using XIP or save amount of flash using compression. By choosing the latter option, the cost saving will generally be greater than the former option , because flash is more expensive than RAM. The operating system is more flexible in its use of the available RAM, discarding file buffers during the periods of high memory pressure. Furthermore , because write operations to flash chips are slower , compressing the data it is faster in most cases.

The main problem with XIP is the interaction with the memory management software. Firstly, for all known memory management units, each page of data must be page-aligned on the flash in order for it to be mapped into processes address space, which makes such a file system to waste. Secondly , while giving write or erase commands to a flash chip, it may return status words on all read cycles, therefore all currently valid mappings would have to be found and invalidated for the duration of operation.

C. Fault Tolerance Support

We can integrate a 32-bit CRC in every block of data, but this only gives error detection - it does not allow the file system to correct errors. We need more sophisticated methods of dealing with single-bit errors in flash chip. Possible solutions are using Hamming code, BCH Code or low-density parity-check codes.

We can use redundancy to improve fault tolerance. Every critical and high used portion of data must have a backup storage. This, of course, will decrease the storage capacity but it can be very useful in case of forced unplugging or system crashes.

D. Use linux kernel CryptoAPI

In the current version we have written our own encryption functions. But the linux kernel provides an extensible way of using cryptography by the means of *CryptoAPI*. The *Scatterlist Crypto API* takes page vectors (scatterlist) as arguments, and works directly on pages. In some cases, this will allow for pages to be encrypted in-place with no copying.

At the lowest level are algorithms, which register dynamically with the API. Transforms are user-instantiated objects, which maintain state, handle all of the implementation logic, manipulating page vectors and provide an abstraction to the underlying algorithms.

E. Transaction Support

For storing database information in AshFS file system, it may be desirable to expose transactions to userspace. Of course, the user space can implement transactions itself, using only the file system functionality required by POSIX but implementing a transaction-based system on top of AshFS would be far less efficient than using the internal functions of the file system.

V. CONCLUSIONS

Our tests have revealed that Linux VFS offers a framework that supports costly operations such as crypting and compression to be implemented as an intermediate layer in a filesystem without affecting performance to a very large scale.

Across different file system implementations, the speed test for the Read operation yielded uniform results. The difference between implementations is more important for the Write operation, where timings differ to a greater extent. Writing timings grow in a linear way in respect to the file size, because flash drives have constant lookup latency and zero spin time.

Using an official library for the crypting algorithm instead of a custom tailored implementation will improve performance considerably. Symmetric crypting algorithms are also the fastest algorithms available for the needed level of security, which can be increased by selecting a proper length for

the used key, between 128, 192 and 256 bits variants. The compression algorithms can be chosen from a larger number of available implementations but using more complex algorithms than zlib will add higher latencies to the Write operation.

Overall, the AshFS project demonstrates how compression and crypting can be combined in order to achieve higher functionality in the operating system. It is also a proof of concept for how a new file system can be integrated with the existing VFS functionality.

REFERENCES

- [1] David Woodhouse, *JFFS: The Journaling Flash File System*, Ottawa Linux Symposium, 2001
- [2] Han-Joon Kim, Sang-Goo Lee, *A new flash memory management for flash storage system*, COMPSAC '99
- [3] Charles Manning, *YAFFS: the NAND-specific flash file system*, Linuxdevices.org, September 20th 2002
- [4] Eran Gal, Sivan Toledo, *A Transactional Flash File System for Micro-controllers*, USENIX '05
- [5] Seung-Ho Lim, Kyu-Ho Park, *An efficient NAND flash file system for flash memory storage*, IEEE Transaction on Computers, 55th Vol., July 2006